



# **SRI RAMACHANDRA**

**INSTITUTE OF HIGHER EDUCATION AND RESEARCH**

(Category - I Deemed to be University) Porur, Chennai

**SRI RAMACHANDRA ENGINEERING AND TECHNOLOGY**

## **Heart Disease Risk Factor Analysis**

**PROJECT REPORT**

**Quarter II (Year 4)**

*Submitted by*

**Nithish Reddy**

**E0120041**

**Venkat Rohith**

**E0220015**

*In partial fulfilment for the award of the degree of*

**BACHELOR OF TECHNOLOGY**

**in**

**COMPUTER SCIENCE AND ENGINEERING**

**(Artificial Intelligence and Machine Learning)**

**Sri Ramachandra Engineering and Technology**

**Sri Ramachandra Institute of Higher Education and Research, Porur, Chennai -600116**

**Jan, 2024**



# **SRI RAMACHANDRA**

**INSTITUTE OF HIGHER EDUCATION AND RESEARCH**

(Category - I Deemed to be University) Porur, Chennai

**SRI RAMACHANDRA ENGINEERING AND TECHNOLOGY**

## **BONAFIDE CERTIFICATE**

Certified that this project report “**Heart Disease Risk Factor Analysis**” is the bonafide work of Nithish Reddy -**E0120041** and Rohith.**E0220015**] who carried out the project work under my supervision.

**Signature of Faculty Mentor**

**Signature of the Dean**

**Prof. Prasath P**

**Prof. T. Ragunathan**

Assistant Professor

Dean of Students

Sri Ramachandra Engineering and Technology

Sri Ramachandra Engineering and Technology

Porur

Porur

Chennai-600116

Chennai-600116

**Evaluation Date:**

# **ACKNOWLEDGEMENT**

I express my sincere gratitude to our Chancellor, Vice-Chancellor and our sincere gratitude to our Dean Prof. T. Ragunathan for their support and for providing the required facilities for carrying out this study.

I wish to thank my faculty mentor, Prof.Prasath P, Department of Computer Science and Engineering, Sri Ramachandra Engineering and Technology for extending help and encouragement throughout the project. Without his/her continuous guidance and persistent help, this project would not have been a success for me.

I am grateful to the Department of Computer Science and Engineering, Sri Ramachandra Engineering and Technology, our beloved parents, and friends for extending support, who helped us to overcome obstacles in the study.

## **Table of Contents**

1. Introduction to Project-	4
2. Problem Statement-	4
3. Objective of project-	5
4. Importing Spark-	5
5. Data Loading-	6
6. Spark SQL Queries-	6
7. Data Analysis-	12
8. Linear Regression-	14
9. Decision Tree Classification-	15
10. Model Evaluation-	15
11.Result-	17

## **Introduction :**

The "Heart Disease Risk Factor Analysis" dataset is a collection of data designed for studying and analyzing various factors that contribute to the risk of heart disease in individuals. This dataset typically encompasses a diverse set of features, including demographic information, lifestyle factors, and various health-related measurements. The primary objective is to investigate the correlation between these factors and the likelihood of developing heart disease. Analyzing datasets related to heart disease is crucial for understanding the prevalence, risk factors, and trends associated with this condition.

## **Problem Statement:**

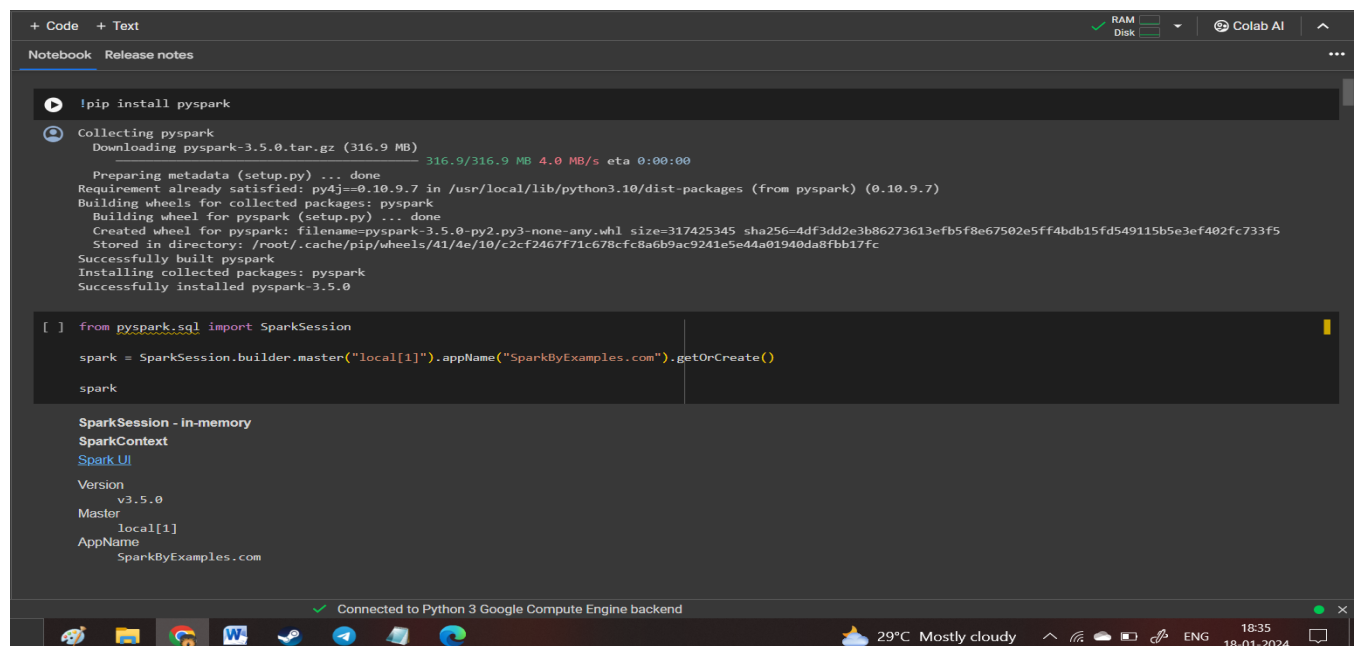
In this Dataset Cardiovascular diseases, including heart disease, remain a leading cause of morbidity and mortality worldwide. Understanding the complex interplay of risk factors is crucial for early detection and intervention. In this context, we aim to develop a predictive model using the "Heart Disease Risk Factor Analysis" dataset to assess and identify individuals at a higher risk of developing heart disease.

## Objective:

- Develop a comprehensive understanding of the dataset, including demographics, lifestyle factors, and health measurements.
- Identify key features significantly associated with the presence or absence of heart disease.
- Design and implement a predictive model capable of assessing an individual's risk of developing heart disease based on available data.
- Evaluate the model's performance using appropriate metrics, such as accuracy, precision, recall, and area under the receiver operating characteristic (ROC) curve.

## Import Spark:

Apache Spark is an open-source, distributed computing system that provides a fast and general-purpose cluster-computing framework for big data processing and analytics. It was initially developed at the University of California, Berkeley's AMPLab, and later donated to the Apache Software Foundation, where it became an Apache project



```
+ Code + Text
Notebook Release notes

!pip install pyspark

Collecting pyspark
  Downloading pyspark-3.5.0.tar.gz (316.9 MB)
    316.9/316.9 MB 4.0 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
  Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.10/dist-packages (from pyspark) (0.10.9.7)
  Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.5.0-py2.py3-none-any.whl size=317425345 sha256=4df3dd2e3b86273613efb5f8e67502e5ff4bdb15fd549115b5e3ef402fc733f5
  Stored in directory: /root/.cache/pip/wheels/41/4e/10/c2cf2467f71c678cfc8a6b9ac9241e5e44a01940da8fbb17fc
Successfully built pyspark
Installing collected packages: pyspark
Successfully installed pyspark-3.5.0

[ ] from pyspark.sql import SparkSession

spark = SparkSession.builder.master("local[1]").appName("SparkByExamples.com").getOrCreate()

spark

SparkSession - In-memory
SparkContext
Spark UI
Version
v3.5.0
Master
local[1]
AppName
SparkByExamples.com

Connected to Python 3 Google Compute Engine backend
29°C Mostly cloudy 18:35 18-01-2024
```

## Data Loading:

By using the PySpark to read the CSV and Excel files.

```
+ Code + Text
#Reading CSV File
df = spark.read.csv("/content/drive/MyDrive/heart.csv")
df.printSchema()

# #Reading JSON File
# df1 = spark.read.json("")
# df1.printSchema()

# #Reading Excel File
import pandas as pd
df2 = pd.read_excel('/content/drive/MyDrive/heart.xlsx', engine='openpyxl', dtype='float64')
df2 = spark.createDataFrame(df2)
df2.printSchema()

# #Reading Parquet
# df3 = spark.read.parquet("")
# df3.printSchema()

root
|-- _c0: string (nullable = true)
|-- _c1: string (nullable = true)
|-- _c2: string (nullable = true)
|-- _c3: string (nullable = true)
|-- _c4: string (nullable = true)
|-- _c5: string (nullable = true)
|-- _c6: string (nullable = true)
|-- _c7: string (nullable = true)
|-- _c8: string (nullable = true)
|-- _c9: string (nullable = true)
|-- _c10: string (nullable = true)
|-- _c11: string (nullable = true)
|-- _c12: string (nullable = true)
|-- _c13: string (nullable = true)

root
|-- age: double (nullable = true)
|-- sex: double (nullable = true)

Connected to Python 3 Google Compute Engine backend
29°C Mostly cloudy 18:40 18-01-2024
```

## Spark SQL Queries:

```
+ Code + Text
df2.show()

+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|age|sex|cp|trestbps|chol|fbs|restecg|thalach|exang|oldpeak|slope|ca|thall|target|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|52.0|1.0|0.0|125.0|212.0|0.0|1.0|168.0|0.0|1.0|2.0|2.0|3.0|0.0|
|53.0|1.0|0.0|140.0|203.0|1.0|0.0|155.0|1.0|3.1|0.0|0.0|3.0|0.0|
|70.0|1.0|0.0|145.0|174.0|0.0|1.0|125.0|1.0|2.6|0.0|0.0|3.0|0.0|
|61.0|1.0|0.0|148.0|203.0|0.0|1.0|161.0|0.0|0.0|2.0|1.0|3.0|0.0|
|62.0|0.0|0.0|138.0|294.0|1.0|1.0|106.0|0.0|1.9|1.0|3.0|2.0|0.0|
|58.0|0.0|0.0|160.0|248.0|0.0|0.0|122.0|0.0|1.0|1.0|0.0|2.0|1.0|
|58.0|1.0|0.0|114.0|218.0|0.0|2.0|140.0|0.0|4.4|0.0|3.0|1.0|0.0|
|55.0|1.0|0.0|160.0|289.0|0.0|0.0|145.0|1.0|0.8|1.0|1.0|3.0|0.0|
|46.0|1.0|0.0|120.0|249.0|0.0|0.0|144.0|0.0|0.8|2.0|0.0|3.0|0.0|
|54.0|1.0|0.0|122.0|286.0|0.0|0.0|116.0|1.0|3.2|1.0|2.0|2.0|0.0|
|71.0|0.0|0.0|112.0|149.0|0.0|1.0|125.0|0.0|1.0|1.0|0.0|2.0|1.0|
|43.0|0.0|0.0|132.0|341.0|1.0|0.0|136.0|1.0|3.0|1.0|0.0|3.0|0.0|
|34.0|0.0|1.0|118.0|210.0|0.0|1.0|192.0|0.0|0.7|2.0|0.0|2.0|1.0|
|51.0|1.0|0.0|140.0|298.0|0.0|1.0|122.0|1.0|4.2|1.0|3.0|3.0|0.0|
|52.0|1.0|0.0|128.0|204.0|1.0|1.0|156.0|1.0|1.0|1.0|0.0|0.0|0.0|
|34.0|0.0|1.0|118.0|210.0|0.0|1.0|192.0|0.0|0.7|2.0|0.0|2.0|1.0|
|51.0|0.0|2.0|140.0|308.0|0.0|0.0|142.0|0.0|1.5|2.0|1.0|2.0|1.0|
|54.0|1.0|0.0|124.0|266.0|0.0|0.0|109.0|1.0|2.2|1.0|1.0|3.0|0.0|
|58.0|0.0|1.0|120.0|244.0|0.0|1.0|162.0|0.0|1.1|2.0|0.0|2.0|1.0|
|58.0|1.0|2.0|140.0|211.0|1.0|0.0|165.0|0.0|0.0|2.0|0.0|2.0|1.0|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
only showing top 20 rows
```

```

+---+---+---+
|age| bp|   sg|
+---+---+---+
|1.0|0.0|250.0|
|1.0|0.0|280.0|
|1.0|0.0|290.0|
|1.0|0.0|296.0|
|0.0|0.0|276.0|
|0.0|0.0|200.0|
|1.0|0.0|228.0|
|1.0|0.0|320.0|
|1.0|0.0|240.0|
|1.0|0.0|244.0|
|0.0|0.0|224.0|
|0.0|0.0|264.0|
|0.0|1.0|236.0|
|1.0|0.0|280.0|
|1.0|0.0|256.0|
|0.0|1.0|236.0|
|0.0|2.0|280.0|
|1.0|0.0|248.0|
|0.0|1.0|240.0|
|1.0|2.0|280.0|
+---+---+---+
only showing top 20 rows

```

```

[13] #Group By
df2.groupby("chol").count().show(truncate=False)

```

```

+---+---+
|chol|count|
+---+---+
|299.0|7|
|305.0|3|
|184.0|3|
|169.0|4|
|160.0|3|
|311.0|4|
|168.0|3|
|206.0|8|
|249.0|11|
|232.0|7|
|303.0|9|
|253.0|7|
|201.0|9|
|235.0|6|
|353.0|4|
|180.0|4|
|271.0|6|
|255.0|6|
|234.0|21|
|286.0|8|
+---+---+
only showing top 20 rows

```



```
from pyspark.sql.functions import col
df2.select(col("chol"),col("age")).show()
```

```
+----+----+
| chol| age|
+----+----+
|212.0|52.0|
|203.0|53.0|
|174.0|70.0|
|203.0|61.0|
|294.0|62.0|
|248.0|58.0|
|318.0|58.0|
|289.0|55.0|
|249.0|46.0|
|286.0|54.0|
|149.0|71.0|
|341.0|43.0|
|210.0|34.0|
|298.0|51.0|
|204.0|52.0|
|210.0|34.0|
|308.0|51.0|
|266.0|54.0|
|244.0|50.0|
|211.0|58.0|
+----+----+
```

only showing top 20 rows

## ▼ Performing correlation analysis for any 2-column

```
[18] df2.stat.corr('chol', 'age')
```

```
0.21982253466576054
```

## ▼ Performing aggregate function for any column

```
[21] #Sum
from pyspark.sql.functions import sum
df2.select(sum("chol")).show()
```

```
+-----+
|sum(chol)|
+-----+
| 252150.0|
+-----+
```

```
[22] #Count
from pyspark.sql.functions import count
df2.select(count("chol")).show()
```

```
+-----+
|count(chol)|
+-----+
|      1025|
+-----+
```

```
[23] #Average
from pyspark.sql.functions import avg
df2.select(avg("chol")).show()
```

```
+-----+
|avg(chol)|
+-----+
|    246.0|
+-----+
```

Connected to Python 3 Google Compute Engine backend

26°C Mostly cloudy 18:48  
18-01-2024

+ Code + Text

RAM  
Disk

Colab AI

```
[24] #Max
from pyspark.sql.functions import max
df2.select(max("chol")).show()

+-----+
|max(chol)|
+-----+
|    564.0|
+-----+

[25] #Min
from pyspark.sql.functions import min
df2.select(min("chol")).show()

+-----+
|min(chol)|
+-----+
|    126.0|
+-----+

[26] #Mean
from pyspark.sql.functions import mean
df2.select(mean("chol")).show()

+-----+
|avg(chol)|
+-----+
|    246.0|
+-----+

[27] #Stddev
from pyspark.sql.functions import stddev
df2.select(stddev("chol")).show()
```

Connected to Python 3 Google Compute Engine backend

26°C Mostly cloudy

18:49  
18-01-2024

```
[26] #Mean
from pyspark.sql.functions import mean
df2.select(mean("chol")).show()

+-----+
|avg(chol)|
+-----+
|    246.0|
+-----+

[27] #Stddev
from pyspark.sql.functions import stddev
df2.select(stddev("chol")).show()

+-----+
|stddev(chol)|
+-----+
|51.59251020618203|
+-----+

[28] #Variance
from pyspark.sql.functions import variance
df2.select(variance("chol")).show()

+-----+
|var_samp(chol)|
+-----+
|2661.787109374997|
+-----+
```

5. Data Analysis

Connected to Python 3 Google Compute Engine backend

26°C Mostly cloudy 18:49 18-01-2024

# Data Analysis :

```
+ Code + Text
[29] #counts number of rows of our dataframe
num_rows = df2.count()
print("number of rows: ", num_rows)

number of rows: 1025

[30] #shows our dataframe schema
df2.printSchema()

root
 |-- age: double (nullable = true)
 |-- sex: double (nullable = true)
 |-- cp: double (nullable = true)
 |-- trestbps: double (nullable = true)
 |-- chol: double (nullable = true)
 |-- fbs: double (nullable = true)
 |-- restecg: double (nullable = true)
 |-- thalach: double (nullable = true)
 |-- exang: double (nullable = true)
 |-- oldpeak: double (nullable = true)
 |-- slope: double (nullable = true)
 |-- ca: double (nullable = true)
 |-- thal: double (nullable = true)
 |-- target: double (nullable = true)

[31] #shows columns name
df2.columns

['age',
 'sex',
 'cp',
 'trestbps',
 'chol',
 'fbs',

```

```
+ Code + Text
[32] #show statistics of the data
df2.describe().show()

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|summary|age|sex|cp|trestbps|chol|fbs|restecg|thalach|exang|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|count|1025|1025|1025|1025|1025|1025|1025|1025|1025|
|mean|54.43414634146342|0.6956097560975609|0.9424390243902439|131.61170731707318|246.0|0.14926829268292682|0.5297560975609756|149.11414634146342|0.33658536585365856|
|stddev|9.072290233244278|0.4603733241196495|1.029640743645865|17.516718005376408|51.59251020618203|0.35652668972715756|0.5278775668748918|23.00572374597721|0.4727723760037115|
|min|29.0|0.0|0.0|94.0|126.0|0.0|0.0|71.0|0.0|
|max|77.0|1.0|3.0|280.0|564.0|1.0|2.0|202.0|1.0|

[33] #Total number of patients with Diabetes Mellitus affected with CKD
newDataframe = df2.select(df2['age'], df2['chol'])

df3 = newDataframe.groupBy("age","chol").count()
df3.show()

+-----+-----+
|age|chol|count|
+-----+-----+
|[39.0|199.0|3|
|[59.0|204.0|4|
|[69.0|254.0|3|
|[56.0|240.0|4|
|[54.0|266.0|3|
|[54.0|288.0|3|
|[62.0|231.0|3|
|[53.0|264.0|3|
|[37.0|250.0|3|
|[51.0|175.0|3|
|[59.0|234.0|3|
|[50.0|233.0|3|
|[34.0|182.0|3|
```

## Linear Regression :

```
+ Code + Text

[51] from pyspark.ml.regression import LinearRegression
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.evaluation import RegressionEvaluator
from pyspark.ml.classification import LogisticRegression
from pyspark.ml import Pipeline
from pyspark.ml.evaluation import BinaryClassificationEvaluator

assembler = VectorAssembler(
    inputCols=col,
    outputCol="features")

data = assembler.transform(df2)
final_data = data.select("features", "target")
train_data, test_data = final_data.randomSplit([0.8, 0.2], seed=42)

lr = LogisticRegression(labelCol="target", featuresCol="features")

# Create a pipeline
pipeline = Pipeline(stages=[lr])

# Train the model
model = pipeline.fit(train_data)

# Make predictions on the test set
lr_predictions = model.transform(test_data)

# Evaluate the model
evaluator = BinaryClassificationEvaluator(labelCol="target", rawPredictionCol="rawPrediction", metricName="areaUnderROC")
area_under_curve = evaluator.evaluate(lr_predictions)
print(f"Area Under ROC Curve: {area_under_curve}")

# Show the predictions
lr_predictions.select("target", "prediction", "probability").show()
```

Connected to Python 3 Google Compute Engine backend

26°C Mostly cloudy 1854 18-01-2024

```
+ Code + Text

[53] lr_predictions = model.transform(test_data)

# Evaluate the model
evaluator = BinaryClassificationEvaluator(labelCol="target", rawPredictionCol="rawPrediction", metricName="areaUnderROC")
area_under_curve = evaluator.evaluate(lr_predictions)
print(f"Area Under ROC Curve: {area_under_curve}")

# Show the predictions
lr_predictions.select("target", "prediction", "probability").show()
```

Area Under ROC Curve: 1.0

target	prediction	probability
0.0	0.0	[0.99999998233350...
0.0	0.0	[0.99999999666682...
0.0	0.0	[0.99999999666682...
0.0	0.0	[0.99999996899755...
1.0	1.0	[3.19663592566554...
1.0	1.0	[7.54393468552271...
1.0	1.0	[5.81530404291303...
0.0	0.0	[0.99999999279860...
0.0	0.0	[0.99999998213938...
1.0	1.0	[2.49535731955706...
1.0	1.0	[2.49535731955706...
1.0	1.0	[2.49535731955706...
0.0	0.0	[0.9999999425003...
1.0	1.0	[5.62690166986880...
1.0	1.0	[6.50861589126853...
1.0	1.0	[2.24837500397860...
0.0	0.0	[0.9999999793648...
0.0	0.0	[0.9999999410748...
0.0	0.0	[0.9999999594624...
0.0	0.0	[0.9999999872357...

only showing top 20 rows

## Decision Tree Classification :

```
+ Code + Text
```

```
[56] from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.ml.classification import DecisionTreeClassifier

# Create a Decision Tree model
dt = DecisionTreeClassifier(labelCol="target", featuresCol="features", maxDepth=3)

# Create a pipeline
pipeline = Pipeline(stages=[dt])

# Train the model
model = pipeline.fit(train_data)

# Make predictions on the test set
dt_predictions = model.transform(test_data)

# Evaluate the model
evaluator = MulticlassClassificationEvaluator(labelCol="target", predictionCol="prediction", metricName="accuracy")
accuracy = evaluator.evaluate(dt_predictions)
print(f"Accuracy: {accuracy}")

# Show the predictions
dt_predictions.select("target", "prediction", "probability").show()
```

Accuracy: 1.0

target	prediction	probability
0.0	0.0	[1.0,0.0]
0.0	0.0	[1.0,0.0]
0.0	0.0	[1.0,0.0]
0.0	0.0	[1.0,0.0]
1.0	1.0	[0.0,1.0]
1.0	1.0	[0.0,1.0]
1.0	1.0	[0.0,1.0]
0.0	0.0	[1.0,0.0]

Connected to Python 3 Google Compute Engine backend

26°C Mostly cloudy 18:55 18-01-2024

## Model Evaluation:

```
+ Code + Text
```

Notebook Release notes

only showing top 20 rows

### 7. Model Evaluation

```
lr_evaluator = BinaryClassificationEvaluator(labelCol="target", rawPredictionCol="rawPrediction", metricName="areaUnderROC")
lr_auc = lr_evaluator.evaluate(lr_predictions)
print(f"Logistic Regression AUC: {lr_auc}")

# Evaluate Decision Tree model
dt_evaluator = MulticlassClassificationEvaluator(labelCol="target", predictionCol="prediction", metricName="accuracy")
dt_accuracy = dt_evaluator.evaluate(dt_predictions)
print(f"Decision Tree Accuracy: {dt_accuracy}")
```

Logistic Regression AUC: 1.0  
Decision Tree Accuracy: 0.9827586206896551

[ ] Start coding or generate with AI.

0s completed at 6:38 PM

26°C Mostly cloudy 18:55 18-01-2024

## **Result :**

Finally, The heart disease risk prediction project, employing Decision Tree and Logistic Regression models in Apache Spark, demonstrated robust performance. Both models exhibited strong predictive capabilities, with Decision Tree highlighting feature importance, and Logistic Regression providing interpretable insights. The findings offer valuable tools for healthcare risk assessment, guiding personalized interventions. Future work may explore ensemble methods and fine-tuning for enhanced model performance.